

```

!wget https://s3.amazonaws.com/keras-file_pathsets/jena_climate_2009_2016.csv.zip
!unzip jena_climate_2009_2016.csv.zip

--2025-04-06 07:36:52-- https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.146.181, 52.217.96.14, 16.182.66.56, ...
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.146.181|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 13565642 (13M) [application/zip]
Saving to: 'jena_climate_2009_2016.csv.zip'

jena_climate_2009_2 100%[=====] 12.94M 18.7MB/s in 0.7s

2025-04-06 07:36:53 (18.7 MB/s) - 'jena_climate_2009_2016.csv.zip' saved [13565642/13565642]

Archive: jena_climate_2009_2016.csv.zip
  inflating: jena_climate_2009_2016.csv
  inflating: __MACOSX/._jena_climate_2009_2016.csv

import os
file_path = os.path.join("jena_climate_2009_2016.csv")

with open(file_path) as f:
    file_path = f.read()

climate_records = file_path.split("\n")
column_headers = climate_records[0].split(",")
climate_records = climate_records[1:]
print(column_headers)
print(len(climate_records))
import os
file_path = os.path.join("jena_climate_2009_2016.csv")

["Date Time", "p (mbar)", "T (degC)", "Tpot (K)", "Tdew (degC)", "rh (%)", "VPmax (mbar)", "VPact (mbar)"
420451

with open(file_path) as f:
    file_path = f.read()

climate_records = file_path.split("\n")
column_headers = climate_records[0].split(",")
climate_records = climate_records[1:]
print(column_headers)
print(len(climate_records))
#Loading and exploring the input data.


["Date Time", "p (mbar)", "T (degC)", "Tpot (K)", "Tdew (degC)", "rh (%)", "VPmax (mbar)", "VPact (mbar)"
420451

import numpy as np
temp = np.zeros((len(climate_records),))
original_file_path = np.zeros((len(climate_records), len(column_headers) - 1))
for i, line in enumerate(climate_records):
    values = [float(x) for x in line.split(",")[1:]]
    temp[i] = values[1]
    original_file_path[i, :] = values[2:]
import numpy as np
temp = np.zeros((len(climate_records),))
original_file_path = np.zeros((len(climate_records), len(column_headers) - 1))
for i, line in enumerate(climate_records):
    values = [float(x) for x in line.split(",")[1:]]
    temp[i] = values[1]
    original_file_path[i, :] = values[2:]
#Plotting the temperature time series from the dataset.

from matplotlib import pyplot as plt
plt.plot(range(len(temp)), temp)
#Plotting temperature trends for the initial 10 days


```

↗ [`matplotlib.lines.Line2D` at `0x7e8095cfdd10`]



```
plt.plot(range(1440), temp[:1440])
# Determining dataset split sizes
```

↗ [`matplotlib.lines.Line2D` at `0x7e809530d490`]



```
num_train_samples = int(0.5 * len(original_file_path))
val_original_file_path = int(0.25 * len(original_file_path))
test_Sample_Number = len(original_file_path) - num_train_samples - val_original_file_path
print("train_Sample_Num:", num_train_samples)
print("val_original_file_path:", val_original_file_path)
print("test_Sample_Number:", test_Sample_Number)
# Dataset preparation
#Normalizing the dataset features
```

↗

```
train_Sample_Num: 210225
val_original_file_path: 105112
test_Sample_Number: 105114
```

```
mean = original_file_path[:num_train_samples].mean(axis=0)
original_file_path -= mean
std = original_file_path[:num_train_samples].std(axis=0)
original_file_path /= std
import numpy as np
from tensorflow import keras
num_series = np.arange(10)
dummy_dataset = keras.utils.timeseries_dataset_from_array(
    data=num_series[:-3],
    targets=num_series[3:],
    sequence_length=3,
    batch_size=2,
)
```

```
for inputs, targets in dummy_dataset:
    for i in range(inputs.shape[0]):
```

```

    print([int(x) for x in inputs[i]], int(targets[i]))
#Creating datasets for training, validation, and testin

↵ [0, 1, 2] 3
   [1, 2, 3] 4
   [2, 3, 4] 5
   [3, 4, 5] 6
   [4, 5, 6] 7

sampling_interval = 6
seq_len = 120
prediction_offset = sampling_interval * (seq_len + 24 - 1)
batch_size = 256

train_file_pathset = keras.utils.timeseries_dataset_from_array(
    original_file_path[:-prediction_offset],
    targets=temp[prediction_offset:],
    sampling_rate=sampling_interval,
    sequence_length=seq_len,
    shuffle=True,
    batch_size=batch_size,
    start_index=0,
    end_index=num_train_samples)

val_file_pathset = keras.utils.timeseries_dataset_from_array(
    original_file_path[:-prediction_offset],
    targets=temp[prediction_offset:],
    sampling_rate=sampling_interval,
    sequence_length=seq_len,
    shuffle=True,
    batch_size=batch_size,
    start_index=num_train_samples,
    end_index=num_train_samples + val_original_file_path)

test_file_pathset = keras.utils.timeseries_dataset_from_array(
    original_file_path[:-prediction_offset],
    targets=temp[prediction_offset:],
    sampling_rate=sampling_interval,
    sequence_length=seq_len,
    shuffle=True,
    batch_size=batch_size,
    start_index=num_train_samples + val_original_file_path)
#Inspecting a sample batch from the dataset to evaluate structure and content

for samples, targets in train_file_pathset:
    print("samples shape:", samples.shape)
    print("targets shape:", targets.shape)
    break
#A common-sense, non-machine-learning baseline
#Measuring baseline model performance using MAE

↵ samples shape: (256, 120, 14)
   targets shape: (256,)

def evaluate_naive_method(file_pathset):
    total_abs_err = 0.
    samples_seen = 0
    for samples, targets in file_pathset:
        preds = samples[:, -1, 1] * std[1] + mean[1]
        total_abs_err += np.sum(np.abs(preds - targets))
        samples_seen += samples.shape[0]
    return total_abs_err / samples_seen

print(f"Validation MAE: {evaluate_naive_method(val_file_pathset):.2f}")
print(f"Test MAE: {evaluate_naive_method(test_file_pathset):.2f}")

↵ Validation MAE: 2.44
   Test MAE: 2.62

!pip install tensorflow==2.12

```

```

Collecting tensorflow==2.12
  Downloading tensorflow-2.12.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (3.4 kB)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (1.6.0)
Requirement already satisfied: flatbuffers>=2.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (25.2.0)
Collecting gast<=0.4.0,>=0.2.1 (from tensorflow==2.12)
  Downloading gast-0.4.0-py3-none-any.whl.metadata (1.1 kB)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.2.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (1.62.0)
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (3.13.0)
Requirement already satisfied: jax>=0.3.15 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.5.2)
Collecting keras<2.13,>=2.12.0 (from tensorflow==2.12)
  Downloading keras-2.12.0-py2.py3-none-any.whl.metadata (1.4 kB)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (18.1.0)
Collecting numpy<1.24,>=1.22 (from tensorflow==2.12)
  Downloading numpy-1.23.5-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.3 kB)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (24.2)
Collecting protobuf!=4.21.0,!<4.21.1,!<4.21.2,!<4.21.3,!<4.21.4,!<4.21.5,<5.0.0dev,>=3.20.3 (from tensorflow==2.12)
  Downloading protobuf-4.25.6-cp37-abi3-manylinux2014_x86_64.whl.metadata (541 bytes)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (1.17.0)
Collecting tensorboard<2.13,>=2.12 (from tensorflow==2.12)
  Downloading tensorboard-2.12.3-py3-none-any.whl.metadata (1.8 kB)
Collecting tensorflow-estimator<2.13,>=2.12.0 (from tensorflow==2.12)
  Downloading tensorflow_estimator-2.12.0-py2.py3-none-any.whl.metadata (1.3 kB)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (3.0.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (4.12.0)
Collecting wrapt<1.15,>=1.11.0 (from tensorflow==2.12)
  Downloading wrapt-1.14.1-cp311-cp311-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2014_x86_64.manylinux2014_x86_64.whl.metadata (3.0 kB)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.37.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.43.0)
Requirement already satisfied: jaxlib<=0.5.2,>=0.5.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.5.2)
Requirement already satisfied: ml_dtypes>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.5.2)
INFO: pip is looking at multiple versions of jax to determine which version is compatible with other requirements. This
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.5.3-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.5.3,>=0.5.3 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.5.3-cp311-cp311-manylinux2014_x86_64.whl.metadata (1.2 kB)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.5.1-py3-none-any.whl.metadata (22 kB)
  Downloading jax-0.5.0-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.5.0,>=0.5.0 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.5.0-cp311-cp311-manylinux2014_x86_64.whl.metadata (978 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.38-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.38,>=0.4.38 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.38-cp311-cp311-manylinux2014_x86_64.whl.metadata (1.0 kB)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.37-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.37,>=0.4.36 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.36-cp311-cp311-manylinux2014_x86_64.whl.metadata (1.0 kB)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.36-py3-none-any.whl.metadata (22 kB)
  Downloading jax-0.4.35-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.35,>=0.4.34 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.35-cp311-cp311-manylinux2014_x86_64.whl.metadata (983 bytes)
INFO: pip is still looking at multiple versions of jax to determine which version is compatible with other requirements.
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.34-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.34,>=0.4.34 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.34-cp311-cp311-manylinux2014_x86_64.whl.metadata (983 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.33-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.33,>=0.4.33 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.33-cp311-cp311-manylinux2014_x86_64.whl.metadata (983 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.31-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.31,>=0.4.30 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.31-cp311-cp311-manylinux2014_x86_64.whl.metadata (983 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.30-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.30,>=0.4.27 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.30-cp311-cp311-manylinux2014_x86_64.whl.metadata (1.0 kB)
Requirement already satisfied: scipy>=1.9 in /usr/local/lib/python3.11/dist-packages (from jax>=0.3.15->tensorflow==2.12) (1.13.0)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (2.35.0)
Collecting google-auth-oauthlib<1.1,>=0.5 (from tensorflow==2.12)
  Downloading google_auth_oauthlib-1.0.0-py2.py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (3.7.0)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (2.32.0)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (0.17.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow==2.12) (3.0.6)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from google-auth<3,>=1.6.3) (5.5.0)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from google-auth<3,>=1.6.3) (0.4.1)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.11/dist-packages (from google-auth<3,>=1.6.3) (4.9.0)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from google-auth-oauthlib<1.1,>=0.5) (2.0.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0) (3.10.1)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0) (2025.1.1)

```

```

Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tenso
Requirement already satisfied: pyasn1<0.7.0,>=0.6.1 in /usr/local/lib/python3.11/dist-packages (from pyasn1-modules>=0.2
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.11/dist-packages (from requests-oauthlib>=0.7.0
Downloading tensorflow-2.12.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (586.0 MB)
586.0/586.0 MB 853.5 kB/s eta 0:00:00
Downloading gast-0.4.0-py3-none-any.whl (9.8 kB)
Downloading jax-0.4.30-py3-none-any.whl (2.0 MB)
2.0/2.0 MB 80.1 MB/s eta 0:00:00
Downloading keras-2.12.0-py2.py3-none-any.whl (1.7 MB)
1.7/1.7 MB 79.4 MB/s eta 0:00:00
Downloading numpy-1.23.5-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.1 MB)
17.1/17.1 MB 101.9 MB/s eta 0:00:00
Downloading protobuf-4.25.6-cp37-abi3-manylinux2014_x86_64.whl (294 kB)
294.6/294.6 kB 27.0 MB/s eta 0:00:00
Downloading tensorboard-2.12.3-py3-none-any.whl (5.6 MB)
5.6/5.6 MB 86.7 MB/s eta 0:00:00
Downloading tensorflow_estimator-2.12.0-py2.py3-none-any.whl (440 kB)
440.7/440.7 kB 35.6 MB/s eta 0:00:00
Downloading wrapt-1.14.1-cp311-cp311-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2_17_x86_64.manylinux2014_x86_64.w
78.4/78.4 kB 8.5 MB/s eta 0:00:00
Downloading google_auth_oauthlib-1.0.0-py2.py3-none-any.whl (18 kB)
Downloading jaxlib-0.4.30-cp311-cp311-manylinux2014_x86_64.whl (79.6 MB)
79.6/79.6 MB 10.0 MB/s eta 0:00:00
Installing collected packages: wrapt, tensorflow-estimator, protobuf, numpy, keras, gast, jaxlib, google-auth-oauthlib,
  Attempting uninstall: wrapt
    Found existing installation: wrapt 1.17.2
    Uninstalling wrapt-1.17.2:
      Successfully uninstalled wrapt-1.17.2
  Attempting uninstall: protobuf
    Found existing installation: protobuf 5.29.4
    Uninstalling protobuf-5.29.4:
      Successfully uninstalled protobuf-5.29.4
  Attempting uninstall: numpy
    Found existing installation: numpy 2.0.2
    Uninstalling numpy-2.0.2:
      Successfully uninstalled numpy-2.0.2
  Attempting uninstall: keras
    Found existing installation: keras 3.8.0
    Uninstalling keras-3.8.0:
      Successfully uninstalled keras-3.8.0
  Attempting uninstall: gast
    Found existing installation: gast 0.6.0
    Uninstalling gast-0.6.0:
      Successfully uninstalled gast-0.6.0
  Attempting uninstall: jaxlib
    Found existing installation: jaxlib 0.5.1
    Uninstalling jaxlib-0.5.1:
      Successfully uninstalled jaxlib-0.5.1
  Attempting uninstall: google-auth-oauthlib
    Found existing installation: google-auth-oauthlib 1.2.1
    Uninstalling google-auth-oauthlib-1.2.1:
      Successfully uninstalled google-auth-oauthlib-1.2.1
  Attempting uninstall: tensorboard
    Found existing installation: tensorboard 2.18.0
    Uninstalling tensorboard-2.18.0:
      Successfully uninstalled tensorboard-2.18.0
  Attempting uninstall: jax
    Found existing installation: jax 0.5.2
    Uninstalling jax-0.5.2:
      Successfully uninstalled jax-0.5.2
  Attempting uninstall: tensorflow
    Found existing installation: tensorflow 2.18.0
    Uninstalling tensorflow-2.18.0:
      Successfully uninstalled tensorflow-2.18.0
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviou
grpcio-status 1.71.0 requires protobuf<6.0dev,>=5.26.1, but you have protobuf 4.25.6 which is incompatible.
tf-keras 2.18.0 requires tensorflow<2.19,>=2.18, but you have tensorflow 2.12.0 which is incompatible.
imbalanced-learn 0.13.0 requires numpy<3,>=1.24.3, but you have numpy 1.23.5 which is incompatible.
chex 0.1.89 requires numpy>=1.24.1, but you have numpy 1.23.5 which is incompatible.
atbucore 0.0.23 requires numpy>=1.24.4, but you have numpy 1.23.5 which is incompatible.
pymc 5.21.1 requires numpy>=1.25.0, but you have numpy 1.23.5 which is incompatible.
blosc2 3.2.1 requires numpy>=1.26, but you have numpy 1.23.5 which is incompatible.
albumintations 2.0.5 requires numpy>=1.24.4, but you have numpy 1.23.5 which is incompatible.
flax 0.10.5 requires jax>=0.5.1, but you have jax 0.4.30 which is incompatible.
treescope 0.1.9 requires numpy>=1.25.2, but you have numpy 1.23.5 which is incompatible.
bigframes 1.42.0 requires numpy>=1.24.0, but you have numpy 1.23.5 which is incompatible.
orbax-checkpoint 0.11.10 requires jax>=0.5.0, but you have jax 0.4.30 which is incompatible.
tensorflow-text 2.18.1 requires tensorflow<2.19,>=2.18.0, but you have tensorflow 2.12.0 which is incompatible.
xarray 2025.1.2 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
scikit-image 0.25.2 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
Successfully installed gast-0.4.0 google-auth-oauthlib-1.0.0 jax-0.4.30 jaxlib-0.4.30 keras-2.12.0 numpy-1.23.5 protobuf
WARNING: The following packages were previously imported in this runtime:
[gast,jax,jaxlib,keras,numpy,tensorflow,wrapt]
You must restart the runtime in order to use newly installed versions.

```

[RESTART SESSION](#)


```

from tensorflow import keras
from tensorflow.keras import layers

inputs = keras.Input(shape=(seq_len, original_file_path.shape[-1]))
x = layers.Flatten()(inputs)
x = layers.Dense(64, activation="relu")(x)
model_outputs = layers.Dense(1)(x)
model = keras.Model(inputs, model_outputs)

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_dense.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_file_pathset,
                    epochs=10,
                    validation_data=val_file_pathset,
                    callbacks=callbacks)

model = keras.models.load_model("jena_dense.keras")
print(f"Test MAE: {model.evaluate(test_file_pathset)[1]:.2f}")

```

```


Epoch 1/10
819/819 [=====] - 42s 51ms/step - loss: 12.5079 - mae: 2.7378 - val_loss: 11.2760 - val_mae: 2.
Epoch 2/10
819/819 [=====] - 41s 49ms/step - loss: 8.3840 - mae: 2.2759 - val_loss: 13.9872 - val_mae: 2.9
Epoch 3/10
819/819 [=====] - 41s 50ms/step - loss: 7.2584 - mae: 2.1174 - val_loss: 11.5287 - val_mae: 2.6
Epoch 4/10
819/819 [=====] - 42s 51ms/step - loss: 6.5942 - mae: 2.0230 - val_loss: 11.7205 - val_mae: 2.7
Epoch 5/10
819/819 [=====] - 49s 60ms/step - loss: 6.1328 - mae: 1.9554 - val_loss: 12.3230 - val_mae: 2.7
Epoch 6/10
819/819 [=====] - 40s 49ms/step - loss: 5.7205 - mae: 1.8883 - val_loss: 10.6712 - val_mae: 2.5
Epoch 7/10
819/819 [=====] - 49s 59ms/step - loss: 5.4281 - mae: 1.8376 - val_loss: 11.6565 - val_mae: 2.6
Epoch 8/10
819/819 [=====] - 40s 49ms/step - loss: 5.1418 - mae: 1.7892 - val_loss: 13.6404 - val_mae: 2.8
Epoch 9/10
819/819 [=====] - 41s 50ms/step - loss: 4.9173 - mae: 1.7530 - val_loss: 13.1373 - val_mae: 2.8
Epoch 10/10
819/819 [=====] - 41s 50ms/step - loss: 4.7305 - mae: 1.7168 - val_loss: 12.0464 - val_mae: 2.7
405/405 [=====] - 13s 31ms/step - loss: 11.6951 - mae: 2.6879
Test MAE: 2.69

```

```

import matplotlib.pyplot as plt
loss = history.history["mae"]
val_loss = history.history["val_mae"]
epochs = range(1, len(loss) + 1)
plt.figure()
plt.plot(epochs, loss, "bo", label="Training MAE")
plt.plot(epochs, val_loss, "b", label="Validation MAE")
plt.title("Training and validation MAE")
plt.legend()
plt.show()
#Let's try a 1D convolutional model
inputs = keras.Input(shape=(seq_len, original_file_path.shape[-1]))
x = layers.Conv1D(8, 24, activation="relu")(inputs)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 12, activation="relu")(x)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 6, activation="relu")(x)
x = layers.GlobalAveragePooling1D()(x)
model_outputs = layers.Dense(1)(x)
model = keras.Model(inputs, model_outputs)

```



```
callbacks = [
    keras.callbacks.ModelCheckpoint("jena_conv.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_file_pathset,
                    epochs=10,
                    validation_data=val_file_pathset,
                    callbacks=callbacks)
```

```
Epoch 1/10
819/819 [=====] - 76s 92ms/step - loss: 21.8567 - mae: 3.6682 - val_loss: 16.1385 - val_mae: 3.
Epoch 2/10
819/819 [=====] - 85s 104ms/step - loss: 15.6008 - mae: 3.1440 - val_loss: 15.1834 - val_mae: 3.
Epoch 3/10
819/819 [=====] - 84s 102ms/step - loss: 14.4357 - mae: 3.0209 - val_loss: 16.3196 - val_mae: 3.
Epoch 4/10
819/819 [=====] - 80s 97ms/step - loss: 13.6544 - mae: 2.9343 - val_loss: 16.5629 - val_mae: 3.
Epoch 5/10
819/819 [=====] - 79s 96ms/step - loss: 13.0269 - mae: 2.8599 - val_loss: 15.0297 - val_mae: 3.
Epoch 6/10
819/819 [=====] - 74s 91ms/step - loss: 12.5610 - mae: 2.8051 - val_loss: 16.9485 - val_mae: 3.
Epoch 7/10
819/819 [=====] - 77s 94ms/step - loss: 12.1834 - mae: 2.7606 - val_loss: 15.5334 - val_mae: 3.
Epoch 8/10
819/819 [=====] - 77s 93ms/step - loss: 11.8229 - mae: 2.7172 - val_loss: 16.1864 - val_mae: 3.
Epoch 9/10
819/819 [=====] - 75s 92ms/step - loss: 11.5371 - mae: 2.6855 - val_loss: 15.8597 - val_mae: 3.
Epoch 10/10
819/819 [=====] - 77s 93ms/step - loss: 11.2517 - mae: 2.6514 - val_loss: 16.0049 - val_mae: 3.
```

```
model = keras.models.load_model("jena_conv.keras")
print(f"Test MAE: {model.evaluate(test_file_pathset)[1]:.2f}")
#Establishing an initial recurrent model as a baseline.
#A simple LSTM-based model
```

```
405/405 [=====] - 17s 40ms/step - loss: 16.0901 - mae: 3.1756
Test MAE: 3.18
```

```
inputs = keras.Input(shape=(seq_len, original_file_path.shape[-1]))
x = layers.LSTM(16)(inputs)
model_outputs = layers.Dense(1)(x)
model = keras.Model(inputs, model_outputs)
```

```
callbacks = [
    keras.callbacks.ModelCheckpoint("jena_lstm.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_file_pathset,
                    epochs=10,
                    validation_data=val_file_pathset,
                    callbacks=callbacks)
```

```
Epoch 1/10
819/819 [=====] - 99s 119ms/step - loss: 37.8345 - mae: 4.4654 - val_loss: 12.1753 - val_mae: 2.
Epoch 2/10
```



```

819/819 [=====] - 95s 115ms/step - loss: 10.9784 - mae: 2.5854 - val_loss: 10.0824 - val_mae: 2.4
Epoch 3/10
819/819 [=====] - 98s 120ms/step - loss: 9.9057 - mae: 2.4542 - val_loss: 9.9743 - val_mae: 2.4
Epoch 4/10
819/819 [=====] - 98s 120ms/step - loss: 9.5329 - mae: 2.4016 - val_loss: 9.9208 - val_mae: 2.4
Epoch 5/10
819/819 [=====] - 96s 117ms/step - loss: 9.2077 - mae: 2.3601 - val_loss: 9.3302 - val_mae: 2.3
Epoch 6/10
819/819 [=====] - 94s 115ms/step - loss: 8.9401 - mae: 2.3262 - val_loss: 9.2844 - val_mae: 2.3
Epoch 7/10
819/819 [=====] - 97s 118ms/step - loss: 8.7727 - mae: 2.3032 - val_loss: 9.4830 - val_mae: 2.4
Epoch 8/10
819/819 [=====] - 115s 140ms/step - loss: 8.5958 - mae: 2.2800 - val_loss: 9.4252 - val_mae: 2.4
Epoch 9/10
819/819 [=====] - 94s 115ms/step - loss: 8.4724 - mae: 2.2593 - val_loss: 9.3826 - val_mae: 2.3
Epoch 10/10
819/819 [=====] - 92s 112ms/step - loss: 8.3591 - mae: 2.2450 - val_loss: 9.5085 - val_mae: 2.3

```

```

model = keras.models.load_model("jena_lstm.keras")
print(f"Test MAE: {model.evaluate(test_file_pathset)[1]:.2f}")
#Understanding recurrent neural networks
#NumPy implementation of a simple RNN

```

```

↻ 405/405 [=====] - 22s 54ms/step - loss: 11.1936 - mae: 2.6500
Test MAE: 2.65

```

```

import numpy as np
timesteps = 100
input_features = 32
output_features = 64
inputs = np.random.random((timesteps, input_features))
state_t = np.zeros((output_features,))
W = np.random.random((output_features, input_features))
U = np.random.random((output_features, output_features))
b = np.random.random((output_features,))
successive_model_outputs = []
for input_t in inputs:
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
    successive_model_outputs.append(output_t)
    state_t = output_t
final_output_sequence = np.stack(successive_model_outputs, axis=0)
#Initializing a basic RNN using Keras
#RNN layer configured for variable-length input sequences

```

```

num_features = 14
inputs = keras.Input(shape=(None, num_features))
model_outputs = layers.SimpleRNN(16)(inputs)
#RNN layer that outputs only the final timestep

```

```

num_features = 14
steps = 120
inputs = keras.Input(shape=(steps, num_features))
model_outputs = layers.SimpleRNN(16, return_sequences=False)(inputs)
print(model_outputs.shape)
#RNN that provides an output for each input time step

```

```

↻ (None, 16)

```

```

num_features = 14
steps = 120
inputs = keras.Input(shape=(steps, num_features))
model_outputs = layers.SimpleRNN(16, return_sequences=True)(inputs)
print(model_outputs.shape)
#Stacking RNN layers

```

```

↻ (None, 120, 16)

```

```

inputs = keras.Input(shape=(steps, num_features))
x = layers.SimpleRNN(16, return_sequences=True)(inputs)
x = layers.SimpleRNN(16, return_sequences=True)(x)
model_outputs = layers.SimpleRNN(16)(x)
#Implementing advanced strategies with recurrent neural networks
# Using recurrent dropout to reduce overfitting in RNNs
#Training and evaluating an LSTM model with dropout

```

```

inputs = keras.Input(shape=(seq_len, original_file_path.shape[-1]))
x = layers.LSTM(32, recurrent_dropout=0.25)(inputs)
x = layers.Dropout(0.5)(x)

```

```

model_outputs = layers.Dense(1)(x)
model = keras.Model(inputs, model_outputs )

callbacks = [
    keras.callbacks.ModelCheckpoint("jena_lstm_dropout.keras",
                                    save_best_only=True)
]
model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
history = model.fit(train_file_pathset,
                    epochs=10,
                    validation_data=val_file_pathset,
                    callbacks=callbacks)
inputs = keras.Input(shape=(seq_len, num_features))
x = layers.LSTM(32, recurrent_dropout=0.2, unroll=True)(inputs)

↩ Epoch 1/10
819/819 [=====] - 205s 247ms/step - loss: 27.8442 - mae: 3.9003 - val_loss: 9.9011 - val_mae: 2
Epoch 2/10
819/819 [=====] - 197s 241ms/step - loss: 14.8935 - mae: 2.9958 - val_loss: 9.6251 - val_mae: 2
Epoch 3/10
819/819 [=====] - 197s 240ms/step - loss: 14.0237 - mae: 2.9047 - val_loss: 9.2187 - val_mae: 2
Epoch 4/10
819/819 [=====] - 197s 240ms/step - loss: 13.4169 - mae: 2.8417 - val_loss: 9.1934 - val_mae: 2
Epoch 5/10
819/819 [=====] - 195s 238ms/step - loss: 12.8983 - mae: 2.7864 - val_loss: 9.6893 - val_mae: 2
Epoch 6/10
819/819 [=====] - 197s 241ms/step - loss: 12.3605 - mae: 2.7273 - val_loss: 9.3112 - val_mae: 2
Epoch 7/10
819/819 [=====] - 197s 240ms/step - loss: 12.0473 - mae: 2.6941 - val_loss: 9.3193 - val_mae: 2
Epoch 8/10
819/819 [=====] - 197s 240ms/step - loss: 11.8512 - mae: 2.6708 - val_loss: 9.3436 - val_mae: 2
Epoch 9/10
819/819 [=====] - 199s 243ms/step - loss: 11.6496 - mae: 2.6467 - val_loss: 9.5948 - val_mae: 2
Epoch 10/10
819/819 [=====] - 198s 242ms/step - loss: 11.4230 - mae: 2.6220 - val_loss: 9.4550 - val_mae: 2

inputs = keras.Input(shape=(seq_len, original_file_path.shape[-1]))
x = layers.GRU(32, recurrent_dropout=0.5, return_sequences=True)(inputs)
x = layers.GRU(32, recurrent_dropout=0.5)(x)
x = layers.Dropout(0.5)(x)
model_outputs = layers.Dense(1)(x)
model = keras.Model(inputs, model_outputs )

```

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.